

```

1 from sklearn.metrics import accuracy_score, classification_report,
  roc_auc_score, roc_curve, confusion_matrix
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4 import pandas as pd
5 from sklearn.model_selection import train_test_split
6 from sklearn.linear_model import LogisticRegression
7 from sklearn.metrics import accuracy_score

```

```

1 df = pd.read_csv('/content/Credit Risk Benchmark Dataset.csv')
2 df.head()

```

	rev_util	age	late_30_59	debt_ratio	monthly_inc	open_credit	late_90	real_estate	late_60_89	dependents	dlq_2yrs
0	0.006999	38	0	0.302150	5440	4	0	1	0	3	0
1	0.704592	63	0	0.471441	8000	9	0	1	0	0	0
2	0.063113	57	0	0.068586	5000	17	0	0	0	0	0
3	0.368397	68	0	0.296273	6250	16	0	2	0	0	0
4	1.000000	34	1	0.000000	3500	0	0	0	0	1	0

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

[+ Code](#) [+ Text](#)

```


1 # Assuming 'predict' is the target variable
2 X = df[['rev_util', 'age', 'late_30_59', 'debt_ratio', 'monthly_inc', 'open_credit', 'late_90', 'real_estate', 'late_60_89', 'dependents']]
3 y = df['dlq_2yrs']

```

```

1 # Handle missing values (replace with mean for numerical features)
2 for col in X.columns:
3     if X[col].dtype in ['int64', 'float64']:
4         X[col].fillna(X[col].mean(), inplace=True)

```

 <ipython-input-15-413c7dcc32be>:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values is a copy.  
For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value, inplace=True)

```

X[col].fillna(X[col].mean(), inplace=True)
<ipython-input-15-413c7dcc32be>:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  

```
X[col].fillna(X[col].mean(), inplace=True)
```

```

1 # Convert to numerical features (if necessary)
2 X = pd.get_dummies(X, columns=['dependents'], dummy_na=True) # Example: One-hot encode 'dependents'
3 # ... handle other categorical features as needed

```

```


1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

```

1 model = LogisticRegression() # Or any suitable model
2 model.fit(X_train, y_train)

```

 /usr/local/lib/python3.11/dist-packages/sklearn/linear\_model/\_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1): STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
 Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  

```
n_iter_i = _check_optimize_result(
```

▼ LogisticRegression ⓘ ?

LogisticRegression()

```

1 y_pred = model.predict(X_test)
2 accuracy = accuracy_score(y_test, y_pred)
3 print(f"Accuracy: {accuracy}")

```

Accuracy: 0.7262937481304218

```

1 # Model Evaluation
2 print(classification_report(y_test, y_pred))
3 print(f"ROC AUC Score: {roc_auc_score(y_test, model.predict_proba(X_test)[: , 1])}")

```

precision recall f1-score support


0	0.67	0.89	0.76	1671
1	0.83	0.56	0.67	1672
accuracy			0.73	3343
macro avg	0.75	0.73	0.72	3343
weighted avg	0.75	0.73	0.72	3343

ROC AUC Score: 0.8114886939889303

```

1 # Confusion Matrix
2 cm = confusion_matrix(y_test, y_pred)
3 plt.figure(figsize=(4, 3))
4 sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
5             xticklabels=['Predicted 0', 'Predicted 1'],
6             yticklabels=['Actual 0', 'Actual 1'])
7 plt.title('Confusion Matrix')
8 plt.xlabel('Predicted')
9 plt.ylabel('Actual')
10 plt.show()

```



```

1 # ROC Curve
2 fpr, tpr, thresholds = roc_curve(y_test, model.predict_proba(X_test)[: , 1])
3 plt.figure(figsize=(4, 3))
4 plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {roc_auc_score(y_test, model.
5 predict_proba(X_test)[: , 1]):.2f})')
6 plt.plot([0, 1], [0, 1], 'k--') # Random classifier line
7 plt.xlabel('False Positive Rate')
8 plt.ylabel('True Positive Rate')
9 plt.title('ROC Curve')
10 plt.legend(loc='lower right')
11 plt.show()

```

